
Observing Process Behaviour - Exercise

Goals

1. Study the behaviour of process execution by monitoring the process structures via the Linux “/proc” directory.
2. Get some experience with command line interface (shell) in Linux, learn some system commands/programs.
3. Try out simple C programming using fork() and exec() commands.

Guideline

Background: The Linux OS provides a convenient way to explore the system properties, as well as properties and states of the currently running processes. Many of the OS structures are visible as files in the pseudo-file system under the directory /proc.

Step 1: Install ubuntu on Virtual machine.

1. Download virtualbox from <https://www.virtualbox.org/>.
 - a. You can choose any virtual machine software you like (e.g. Parallels Desktop,)
2. Download 32-bit ubuntu from <http://releases.ubuntu.com/16.04/ubuntu-16.04.3-desktop-i386.iso>
 - a. NB: To run the executable files provided, you have to use **32-bit** version.
3. Install ubuntu on virtualbox

Tips: to share files between the host and virtual machines, check the following links:

<https://www.youtube.com/watch?v=ddExu55cJOI> (Windows)

<https://www.youtube.com/watch?v=TcrfrVNNGMU> (Mac OS)

Step 2: Open a terminal and enter command `ls /proc` (any text in **this font** represents commands to be executed). The content of the /proc directory is listed. You will see many number entries (representing directories containing information about processes, for example directory 23406 contains information about process with PID 23406), as well as some more meaningful entries (for example ‘version’ and ‘cpuinfo’).

Note: To open a terminal on ubuntu, right click the desktop and click “Open Terminal”

Step 3: Enter command `cat /proc/version`. This should display the content of the file ‘/proc/version’ that contains information about the OS version. Enter command `cat /proc/cpuinfo` and have a look. You might want to explore the content of other files as well.

Step 4: Enter command `ps`. This will list the processes you have launched. At least two processes will be listed: `'ps'` (the one you just launched and which is producing this output) and a shell process (probably `'bash'`, but could be different, depending on your settings). Write down the PID of the shell process.

Step 5: Enter command `cat /proc/XXXX/stat`, where XXXX is the PID of your shell process from the previous step. You will see a bunch of numbers containing lots of information about this process. You can learn the meaning of the numbers by entering `man proc` command (which would show you the manual for `proc`) and scrolling down. You can see most of this information in human readable form in the file `'/proc/XXXX/status'`.

Step 6: We want to explore the state changes of processes, but there is not much interesting activity by the shell process. In order to see more interesting behaviour, we will launch and observe other, specially created programs. Copy the provided file `'lab1.tar'` into your working directory on the Linux system and untar it by executing command `tar -xvf lab1.tar`. This should extract the following files:

1. *procmon* This is a program which periodically (every second) reads the file `'/proc/PID/stat'` and extracts and prints the process state and the number of jiffies (1/100 of a second) this process spent in kernel and user mode, respectively. This program is launched with one parameter – the PID of the process it has to monitor. *procmon* terminates when it cannot open the `/proc/PID/stat` file – usually because the process PID has terminated.
2. Program files *calclloop* and *cploop*. These are two programs with quite different behaviour:
 - a. The program *calclloop* does the following:
 - i. Goes into a loop (10 iterations) that sleeps for 3 seconds, and then starts another loop that increments a variable 400,000,000 times. The calculation takes approximately 2 secs of real time (this may vary according to load on the system).
 - b. The program *cploop* does the following:
 - i. Creates a file that is 500,000 bytes long (fromfile)
 - ii. Goes into a loop (10 iterations) that sleeps for 3 seconds, and then copies the fromfile to the tofile. The copy operation takes approximately 2 secs of real time (this may vary according to load on the system).
 - iii. The program uses two system calls to copy bytes one at a time.
3. Shell program files *tstcalc* and *tstcp*. *tstcalc* launches the *calclloop* program, gets its PID and then launches *procmon* with this PID to monitor the behaviour of *calclloop*. After the *calclloop* has run for 20s, *tstcalc* terminates it and if *procmon* did not terminate by itself within 1s, it is terminated as well. The output of *procmon* is redirected to file `'calc.log'`, where it can be later examined. *tstcp* does the same for *cploop*, saving the *procmon*'s output in file `'cp.log'` To run these shell programs, simply type in their name, i.e. `tstcalc` and `tstcp`.

- a. To launch an executable file in unix OS, use `cd` to navigate to the folder and execute command `./FileName`, i.e. `./tstcalc` and `./tsttcp`
- b. If you get some error like:
 - i. `procmon: not found` and/or
 - ii. `calclloop: not found`
 run `PATH=./:$PATH` (which adds the current dir to the system variable `$PATH`) first
4. A skeleton of a C program `mon.c`. You will later on fill-in the missed code.
5. A directory that contains the C programs `procmon.c`, `calclloop.c` and `cploop.c` for your perusal.

Step 7: Launch `tstcalc`. Do the same with `tsttcp`.

Step 8: Examine the contents of the files '`calc.log`' and '`cp.log`' and answer the following questions:

1. Explain the changes you see in the state, system time, and user time, for the `calclloop` process relating these changes to the operations completed by `calclloop`.
2. Explain the changes you see in the state, system time, and user time, for the `cploop` process relating these changes to the operations completed by `cploop`.
3. Explain the difference of time spent in system time and user time between the two programs (i.e why would one program spend more time in user mode and/or system mode than the other).

Step 9: OK, all of this was easy. Now let's try to do some programming, playing around with the `fork()` and `exec()` calls we learned about in the lectures. Your goal: write a C program `mon.c` which will provide most of the functionality of `tstcalc` and `tsttcp`:

1. `mon.c` has one command line argument: the name `M` of the program it has to launch
2. `mon.c` will launch the program `M` (the program `M` does not take any parameters) and learn its PID.
3. `mon.c` will launch `procmon` with the argument PID
4. `mon.c` will sleep for 20s, then it will terminate program `M`, will sleep 2 more seconds and will (try to) terminate `procmon`.
5. Do not worry about redirecting the output of `procmon` into a log file (no need to bother you with file I/O in C at this moment)

Step 10: Compile your C program by entering `gcc mon.c -o mon`. If the compilation succeeded, an executable called `mon` has been created.

Step 11: Enter `mon calclloop` and observe. If everything worked well, you should see the same output as you have seen in the log file '`calc.log`'. You can redirect this output into a file by entering `mon calclloop > filename`.

Background you might need:

- The execution of a C program starts in procedure `main()` which has two arguments: integer `argc` containing the number of command line arguments and an `argc`-element array of strings `argv`, containing the command line arguments. By convention, `argv[0]` is the name of the program, `argv[1]` is its first argument.
- the `fork()` command creates a new process. Both the parent and the child continue as if they returned from a `fork()` call, however the parent gets as return value the PID of the child, while the child gets 0. Type **man fork** to read more about `fork()`.
- The `execl(path, arg1, ...)` command replaces the current process with the specified program launched with the provided arguments. Type **man execl** to learn about the exact meaning of its arguments.
The code `'execl("calcloop", "calcloop", NULL)'` will replace the current process with the *calcloop* program. The code `'execl("/bin/ls", "ls", "-l", NULL)'` will launch `ls -l`.
- You will need to convert the integer PID into a string to pass to `procmon`. In C you can do that by calling library function `string *itoa(int num, string *buf, int radix)`. (Check whether this really works on solaris!)
- Function `sleep()` will cause the program to sleep for the specified number of seconds.
- Function `kill(pid, sig)` will send a signal `sig` to process `pid`.
 - have a look at `signal.h` to see the different signals
 - google is your friend, this search result on 'signal.h' might be helpful (there are many more)

<http://www.opengroup.org/onlinepubs/009695399/basedefs/signal.h.html>